

# Week 2

---

## [Week 2](#)

### [Multivariate Linear Regression](#)

#### [Cost function](#)

#### [Gradient Descent for Multiple Variables](#)

#### [Matrix Notation](#)

#### [Algorithm: Gradient Descent](#)

#### [Feature Normalization - Speeding up Gradient Descent](#)

##### [Feature Scaling](#)

##### [Mean Normalization](#)

#### [Learning Rate & Debugging](#)

##### [Debugging gradient descent.](#)

##### [Automatic convergence test.](#)

##### [Example 1: Different Learning Rates](#)

##### [Example 2: Oscillation](#)

#### [Vectorized: Gradient Descent](#)

##### [Derivation](#)

### [Features](#)

### [Polynomial Regression](#)

### [Normal Equation - Computing Parameters Analytically](#)

#### [Formula: Normal Equation](#)

#### [Terminology:](#)

##### [Example:](#)

##### [Octave Code](#)

##### [Intuition](#)

##### [Pros & Cons](#)

##### [Recap:](#)

##### [Non-Inevitability](#)

##### [Common Causes](#)

# Multivariate Linear Regression

---

Linear regression with multiple variables is also known as "multivariate linear regression".

We now introduce notation for equations where we can have any number of input variables.

$x_j^{(i)}$  = value of feature  $j$  in the  $i^{th}$  training example

$x^{(i)}$  = the column vector of all the feature inputs of the  $i^{th}$  training example

$m$  = the number of training examples

$n = |x^{(i)}|$ ; (the number of features)

Now define the multivariable form of the hypothesis function as follows, accommodating these multiple features:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \cdots + \theta_n x_n$$

In order to develop intuition about this function, we can think about  $\theta_0$  as the basic price of a house,  $\theta_1$  as the price per square meter,  $\theta_2$  as the price per floor, etc.  $x_1$  will be the number of square meters in the house,  $x_2$  the number of floors, etc.

Using the definition of matrix multiplication, our multivariable hypothesis function can be concisely represented as:

$$h_{\theta}(x) = [\theta_0 \quad \theta_1 \quad \dots \quad \theta_n] \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} = \theta^T x$$

This is a vectorization of our hypothesis function for one training example; see the lessons on vectorization to learn more.

Remark: Note that for convenience reasons in this course Mr. Ng assumes  $x_0(i) = 1$  for  $(i \in 1, \dots, m)$

[**Note:** So that we can do matrix operations with theta and x, we will set  $x_0^{(i)} = 1$ , for all values of  $i$ . This makes the two vectors  $\theta$  and  $x(i)$  match each other element-wise (that is, have the same number of elements:  $n+1$ ).]

The training examples are stored in X row-wise, like such:

$$X = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} \\ x_0^{(2)} & x_1^{(2)} \\ x_0^{(3)} & x_1^{(3)} \end{bmatrix}, \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix}$$

You can calculate the hypothesis as a column vector of size  $(m \times 1)$  with:

$$h_{\theta}(X) = X\theta$$

**For the rest of these notes, and other lecture notes, X will represent a matrix of training examples  $x(i)$  stored row-wise.**

## Cost function

---

For the parameter vector  $\theta$  (of type  $R^{n+1}$  or in  $R^{(n+1) \times 1}$ , the cost function is:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

The vectorized version is:

$$J(\theta) = \frac{1}{2m} (X\theta - \vec{y})^T (X\theta - \vec{y})$$

Where  $\vec{y}$  denotes the vector of all y values.

## Gradient Descent for Multiple Variables

---

The gradient descent equation itself is generally the same form; we just have to repeat it for our 'n' features:

```

repeat until convergence: {
 $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0^{(i)}$ 
 $\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_1^{(i)}$ 
 $\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_2^{(i)}$ 
...
}

```

In other words:

```

repeat until convergence: {
 $\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$     for j := 0..n
}

```

## Matrix Notation

---

The Gradient Descent rule can be expressed as:

$$\theta := \theta - \alpha \nabla J(\theta)$$

Where  $\nabla J(\theta)$  is a column vector of the form:

$$\nabla J(\theta) = \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_0} \\ \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{bmatrix}$$

The j-th component of the gradient is the summation of the product of two terms:

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_j} &= \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \\ &= \frac{1}{m} \sum_{i=1}^m x_j^{(i)} \cdot (h_{\theta}(x^{(i)}) - y^{(i)}) \end{aligned}$$

Sometimes, the summation of the product of two terms can be expressed as the product of two vectors.

Here,  $x_j^{(i)}$  for  $i = 1, \dots, m$ , represents the m elements of the j-th column,  $\vec{x}_j$ , of the training set X.

The other term  $(h_{\theta}(x^{(i)}) - y^{(i)})$  is the vector of the deviations between the predictions  $h_{\theta}(x^{(i)})$  and the true values  $y^{(i)}$ . Re-writing  $\frac{\partial J(\theta)}{\partial \theta_j}$ , we have:

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \vec{x}_j^T (X\theta - \vec{y})$$

$$\nabla J(\theta) = \frac{1}{m} X^T (X\theta - \vec{y})$$

Finally, the matrix notation (vectorized) of the Gradient Descent rule is:

$$\theta := \theta - \frac{\alpha}{m} X^T (X\theta - \vec{y})$$

## Algorithm: Gradient Descent

Repeat until convergence:

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \text{ for } j := 0 \dots n$$

Context of gradient descent w.r.t.  $n=1$

### Gradient Descent

Previously ( $n=1$ ):

Repeat {

$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$

$\frac{\partial}{\partial \theta_0} J(\theta)$

$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$

(simultaneously update  $\theta_0, \theta_1$ )

}

**New algorithm ( $n \geq 1$ ):**

Repeat {

$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$

(simultaneously update  $\theta_j$  for  $j = 0, \dots, n$ )

}

---

$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$

$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$

$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$

...

*Handwritten notes in red:*

- For  $n=1$ , the first update is for  $\theta_0$  and the second for  $\theta_1$ .
- For  $n \geq 1$ , the updates are for  $\theta_0, \theta_1, \theta_2, \dots$ .
- Handwritten:  $\frac{\partial}{\partial \theta_j} J(\theta)$  with an arrow pointing to the update for  $\theta_j$ .
- Handwritten:  $x_0^{(i)} = 1$  with an arrow pointing to the  $x_0^{(i)}$  term in the first update of the new algorithm.

**Note**  $x_0^i = 1$  and the partial derivative of the cost function for  $n \geq 1$

## Feature Normalization - Speeding up Gradient Descent

We can speed up our algorithm by having our input values in roughly the same ranges.  $\theta$  will descent quickly on small ranges and slowly on large ranges. i.e oscillation occurs on the uneven surface.

Two techniques to help with this are **feature scaling** and **mean normalization**. Feature scaling involves dividing the input values by the range (i.e. the maximum value minus the minimum value) of the input variable, resulting in a new range of just 1. *Mean normalization* involves subtracting the average value for an input variable from the values for that input variable, resulting in a new average value for the input variable of just zero.


## Feature Scaling

Make sure features are on a **similar scale**

Get every feature into approximately a  $a \leq x_i \leq b$  range

Naïve Method:

$$x_i := \frac{x_i}{\max(x_i)}$$



## Mean Normalization

Replace  $x_i$  with  $x_i - \mu_i$  to make features have approximately **zero mean**

**NOTE** do not apply to  $x_0 = 1$

Method: - This method will normalize and feature scale.

$$x_i := \frac{x_i - \mu_i}{s_i}$$

$\mu_i$  — average of all values of feature i

$s_i$  — range of values (min-max) or its the standard deviation

**NOTE** The quizzes in this course use range - the programming exercises use standard deviation.

## Learning Rate & Debugging

- If  $\alpha$  is too small: slow convergence
- If  $\alpha$  is too large:  $J(\theta)$  may not decrease on every iteration; may not converge


### Debugging gradient descent.

Make a plot with *number of iterations* on the x-axis. Now plot the cost function,  $J(\theta)$  over the number of iterations of gradient descent. If  $J(\theta)$  ever increases, then you probably need to decrease  $\alpha$ .

## Automatic convergence test.

Declare convergence if  $J(\theta)$  decreases by less than  $E$  in one iteration, where  $E$  is some small value such as  $10^{-3}$ . However in practice it's difficult to choose this threshold value.

### Example 1: Different Learning Rates




The learning rates:  $\alpha_B \ll \alpha_A \ll \alpha_C$

In graph C, the cost function is increasing- the learning rate is too high. Graph A,B converge to an optimum of the cost function, but graph B does so much slower due to how small  $\alpha$  is.

### Example 2: Oscillation

Observe what the cost function plot looks like when overshooting occurs:



## Vectorized: Gradient Descent

$$\theta := \theta - \frac{\alpha}{m} X^T (X\theta - \vec{y})$$

## Derivation

The Gradient Descent rule can be expressed as:

$$\theta := \theta - \alpha \nabla J(\theta)$$

Where  $\nabla J(\theta)$  is a column vector of the form:

$$\nabla J(\theta) = \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_0} \\ \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{bmatrix}$$

The j-th component of the gradient is the summation of the product of two terms:

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_j} &= \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) \cdot x_j^{(i)} \\ &= \frac{1}{m} \sum_{i=1}^m x_j^{(i)} \cdot \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) \end{aligned}$$

Sometimes, the summation of the product of two terms can be expressed as the product of two vectors.

Here,  $x_j^{(i)}$  for  $i = 1, \dots, m$ , represents the m elements of the j-th column,  $\vec{x}_j$ , of the training set X.

The other term  $(h_{\theta}(x^{(i)}) - y^{(i)})$  is the vector of the deviations between the predictions  $h_{\theta}(x^{(i)})$  and the true values  $y^{(i)}$ . Re-writing  $\frac{\partial J(\theta)}{\partial \theta_j}$ , we have:

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \vec{x}_j^T (X\theta - \vec{y})$$

$$\nabla J(\theta) = \frac{1}{m} X^T (X\theta - \vec{y})$$

Finally, the matrix notation (vectorized) of the Gradient Descent rule is:

$$\theta := \theta - \frac{\alpha}{m} X^T (X\theta - \vec{y})$$

## Features

We can **combine** multiple features into one.

**For example:**

We can combine  $x_1$  and  $x_2$  into a new feature  $x_3 = x_1 x_2$  i.e frontage x depth, no need for two separate parameters.

### Housing prices prediction

$$h_{\theta}(x) = \theta_0 + \theta_1 \times \text{frontage} + \theta_2 \times \text{depth}$$



# Polynomial Regression

---

We can **change the behavior or curve** of our hypothesis function by making it a quadratic, cubic or square root function (or any other form).

If our hypothesis functions is  $h_{\theta}(x) = \theta_0 + \theta_1 x_1$ , we can create additional features based on  $x_1$ .

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_1^2$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^3$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 \sqrt{x_1}$$

**NOTE** Scaling becomes very important when dealing with polynomials terms

eg if  $x_1$  has range 1 - 1000 then range of  $x_1^2$  becomes 1 - 1000000 and that of  $x_1^3$  becomes 1 - 1000000000

**NOTE:** We will look algorithms later that can determine what kind of polynomial to use

## Normal Equation - Computing Parameters Analytically

---

Method for solving for  $\theta$  analytically

**NOTE** Feature scaling is not needed

### Formula: Normal Equation

$$\theta = (X^T X)^{-1} X^T y$$

Where  $\theta$  is a vector containing the parameter values

### Terminology:

- $X$  - design matrix

### Example:

---



Examples:  $m = 4$ .

	Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
$x_0$	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

$m \times (n+1)$

$$y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

$m$ -dimensional vector

$$\theta = (X^T X)^{-1} X^T y$$

## Octave Code

```
pinv(x'*x)*x'*y
```

## Intuition


Here we consider the **1D** case. essentially we find the 1st derivative of the cost function and equate it to zero - as is done to find the optimal solution of the curve.

Intuition: If 1D ( $\theta \in \mathbb{R}$ )

$$\rightarrow J(\theta) = a\theta^2 + b\theta + c$$

$$\frac{d}{d\theta} J(\theta) = \dots \stackrel{\text{set}}{=} 0$$

Solve for  $\theta$



$n$  dimensional extension

$$\theta \in \mathbb{R}^{n+1} \quad J(\theta_0, \theta_1, \dots, \theta_m) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \dots = 0 \quad (\text{for every } j)$$

Solve for  $\theta_0, \theta_1, \dots, \theta_n$

## Pros & Cons

Gradient Descent	Normal Equation
Need to choose alpha	No need to choose alpha
Needs many iterations	No need to iterate
$O(kn^2)$	$O(n^3)$ , need to calculate inverse of $X^T X$
Works well when $n$ is large	Slow if $n$ is very large

With the normal equation, computing the inversion has complexity  $O(n^3)$ . So if we have a very large number of features, the normal equation will be slow. In practice, when  $n$  exceeds 10,000 it might be a good time to go from a normal solution to an iterative process.

## Recap:

Matrix identity

$(X^T X)^{-1}$  is inverse of matrix  $X^T X$ .

Set  $A = X^T X$

$$\boxed{(X^T X)^{-1}} = A^{-1}$$

## Non-Inevitability

What if  $X^T X$  is non-invertible?

**NOTE** Octave already considers this situation with `pinv`

## Common Causes

- Redundant features, where two features are very closely related (i.e. they are linearly dependent)
- Too many features (e.g.  $m \leq n$ ). In this case, delete some features or use "regularization" (to be explained in a later lesson).

Solutions to the above problems include deleting a feature that is linearly dependent with another or deleting one or more features when there are too many features.